



kubernetes


Kubernetes Training

Travis Dent

Certified Kubernetes Application Developer (CKAD)

- This training is based on the exam program for CKAD
- <https://www.cncf.io/certification/ckad/>

Logistics

- Clone the Github Repo
- You'll need a text editor and a terminal (for Windows, PowerShell)
- You'll be using Google Kubernetes Engine
- Topic exercises and a use case
- This is a very practical training 

Let's take a look at the repo! 🙄

Setting up access to GKE

You need a Google Cloud Platform account.

Locally you need:

- gcloud
- kubectl

Instructions in the repo `README.md`

Setting up access to GKE

Tooling setup is detailed in training repo `README.md`

Clusters have been created in a GCP Project "[Insert project name]", project ID [Insert project ID].

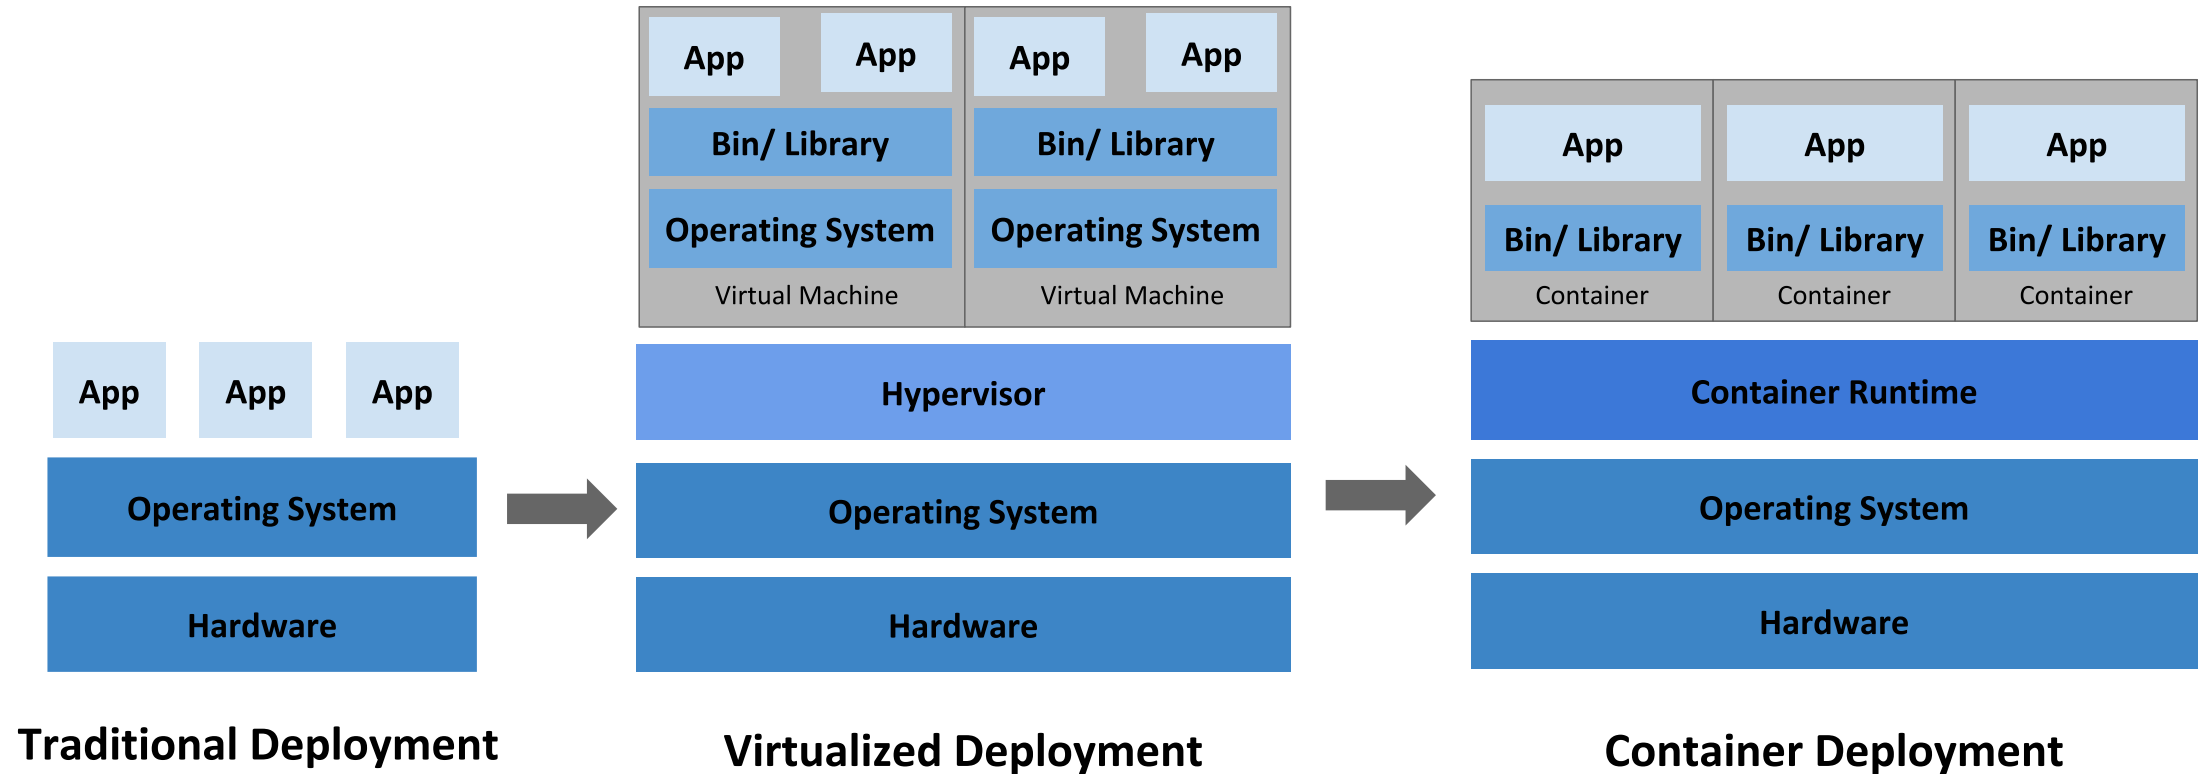
Your cluster will be called `training-<your name>`.

Let's take some time to ensure everyone has access!

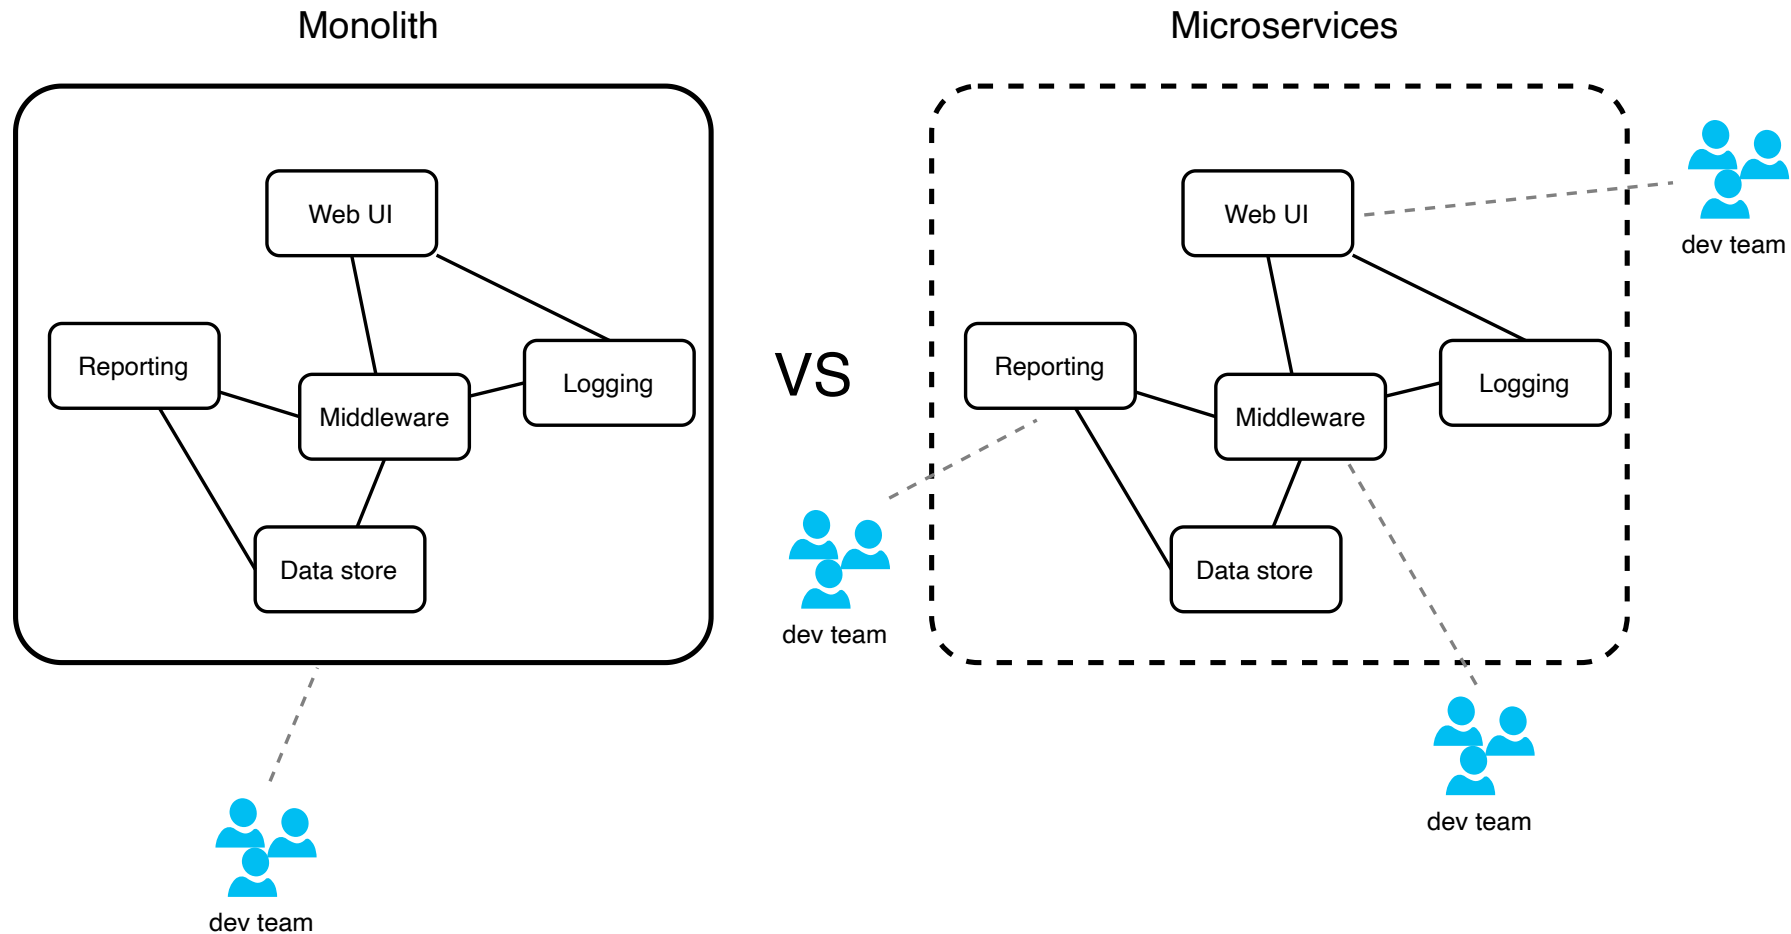
What would you like to learn?

Introduction and theory

Containers FTW



Monoliths vs Microservices



Some history

- Selling servers, network switches, monolith app licenses = big 💰
- 2006: Amazon launches AWS - the birth of cloud computing! 🌟
- 2013: Docker initial release
- 2014: Google introduces Kubernetes
- 2015: Kubernetes donated to Cloud Native Computing Foundation
- 2018: Kubernetes GA on AWS (also Azure, Digital Ocean, more..)

Some more history

- Google used Linux containers to run servers
- Billions per week
- Used a technology called *Borg* and then *Omega*
- Same engineers built *Kubernetes*
- *Kubernetes* based on κυβερνήτης: helmsman

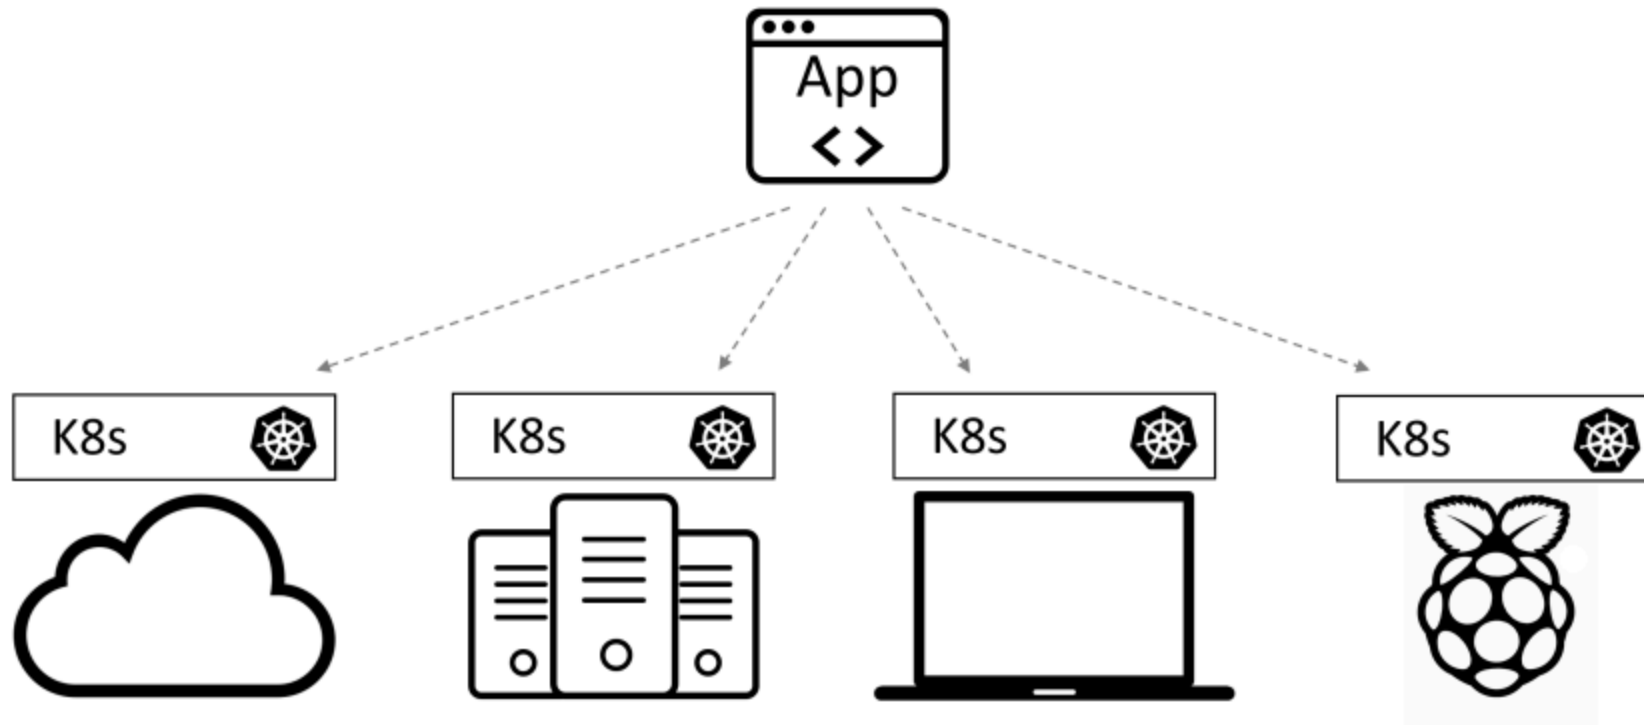


Kubernetes

"Kubernetes is an open-source container orchestration system for automating the deployment, scaling and management of containerized applications."

What makes K8s so attractive?

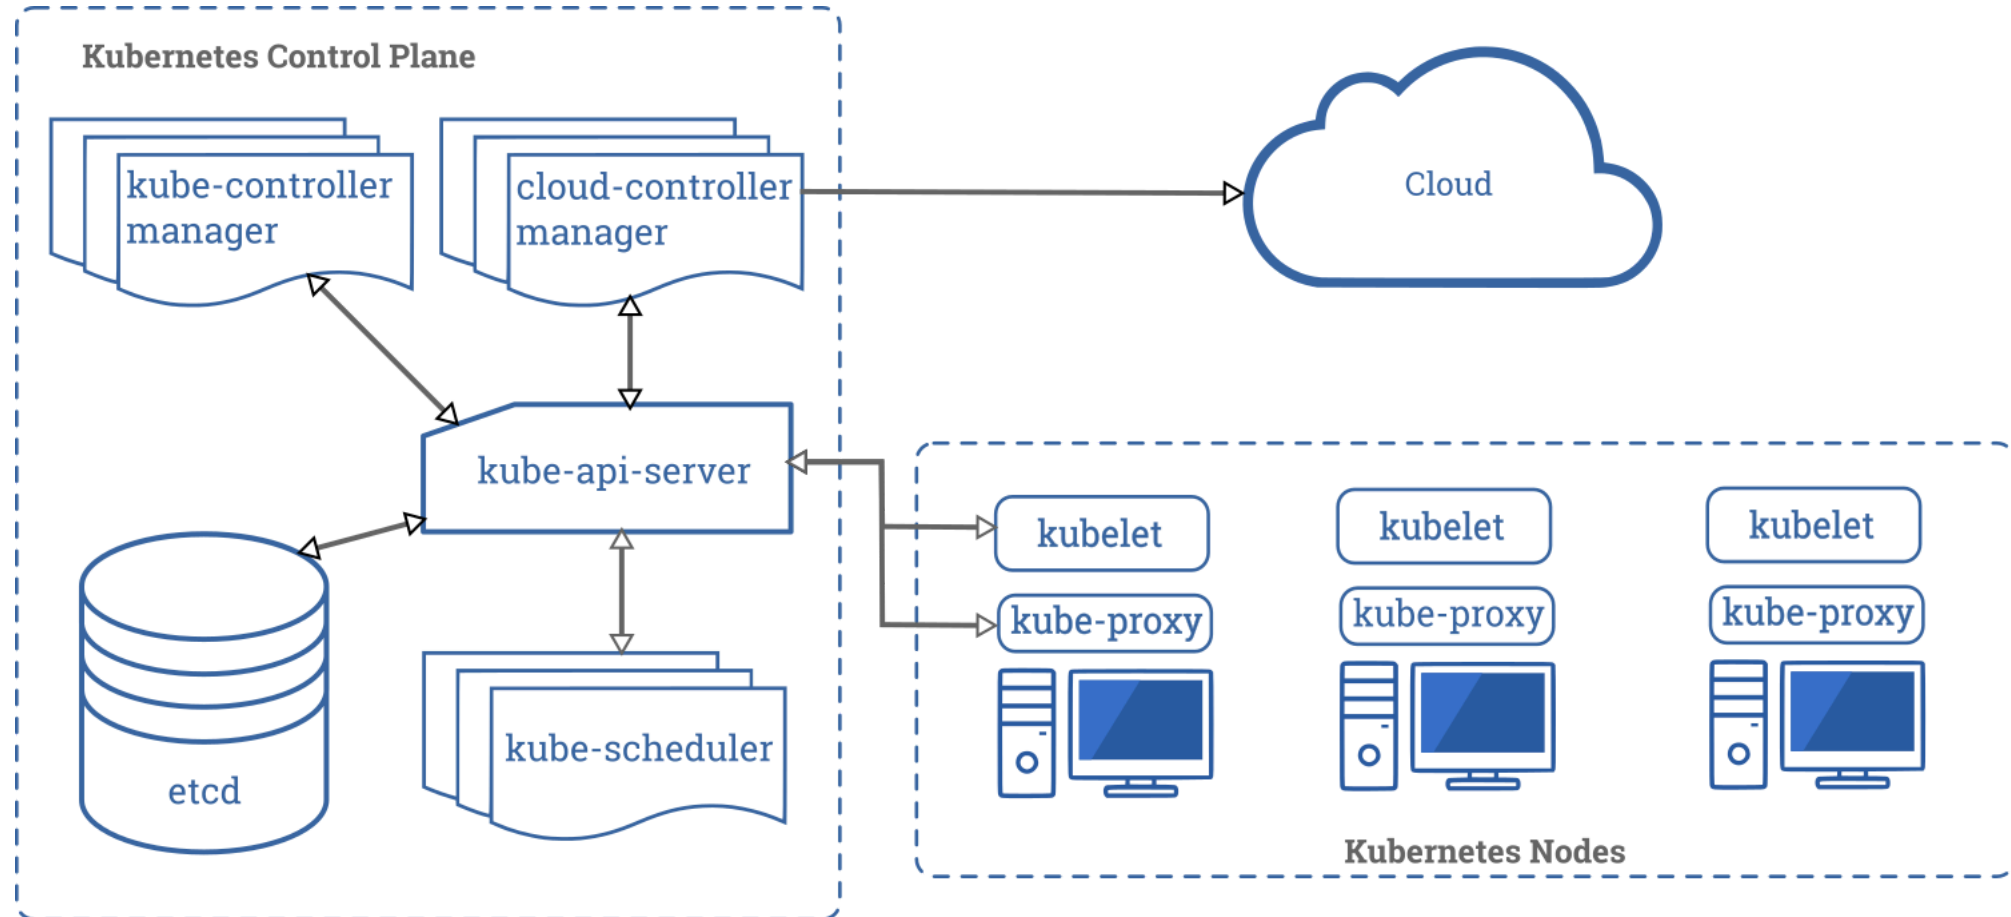
The OS of the cloud



Kubernetes

- Written in Go
- Service discovery and load balancing
- Storage orchestration
- Automated rollouts and rollbacks
- Best use of resources
- Self healing
- Secret and configuration management
- Horizontal scaling
- IPv4/IPv6 dual-stack
- Designed for extensibility

Under the hood



kubectl

- Talks to kube-api-server
- CLI for managing entire clusters
- Check out the kubectl cheatsheet ([link in notes](#))

kubectl [command] [type] [name] [flags]

- Command [get | create | describe]
- Type: [pod | namespace | node]
- Name: name of the resource
- Flags: optional flags
- <https://kubernetes.io/docs/reference/kubectl/#operations>
- <https://kubernetes.io/docs/reference/kubectl/#resource-types>

Kubernetes Objects

- Every interaction with kubectl works on objects
- All "things" in Kubernetes are objects in the kube-api-server:
 - Pods
 - Namespaces
 - ConfigMaps
 - Nodes
 - RBAC
 - Jobs

Kubernetes manifests

- A pod manifest:

```
apiVersion: v1 # schema version
kind: Pod # what kind of object, mapped to REST resource
metadata: # object's metadata
  name: my-pod
  namespace: my-namespace
spec: # Specification of desired behaviour
  containers:
    - name: nginx
      image: nginx:stable
```

- YAML or JSON
- Declarative way to define desired state
- Multiple objects possible (separate by `---`)

Kubernetes Documentation ❤️

Remember this one? **Sunny Bikes**

Sunny bikes is a bike sharing company. Currently sharing bikes in Austin (Texas), New York and San Francisco.

Their system consists of several components written in Python, but they're facing deployment issues because all components work with different versions and have different dependencies.

Sunny Bikes goes Kubernetes!

During the training the Sunny Bikes solution will be transformed into a Kubernetes implementation.

Namespaces

Namespace

- "group" for objects
- Default is `default`
- Allows resource quotas and access isolation
- Not all objects are "namespaced", try `kubectl api-resources`

```
apiVersion: v1
kind: Namespace
metadata:
  name: my-namespace
```

Pods


Nodes

- virtual or physical machine
- contains services for running pods




Pods

- basic unit of Kubernetes
- keep container(s) running
- shared storage and network
- usually one main application
- each pod has own IP



Pods run on Nodes



Pod

- Create yaml file "pod.yaml" (or any other name)
- Deploy to Kubernetes with "kubectl apply -f pod.yaml"

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
  - name: myapp-container
    image: busybox:stable
    command: ['sh', '-c', 'echo Hello Kubernetes! && sleep 3600']
```

Container configuration


- `command` and `args`
- Ports
- `restartPolicy`

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  restartPolicy: Never
  containers:
  - name: myapp-container
    image: busybox:stable
    command: ["/bin/sh"]
    args: ["-c", "while true; do echo hello; sleep 10;done"]
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  restartPolicy: Always
  containers:
  - name: myapp-container
    image: busybox:stable
    ports:
      - containerPort: 80
```

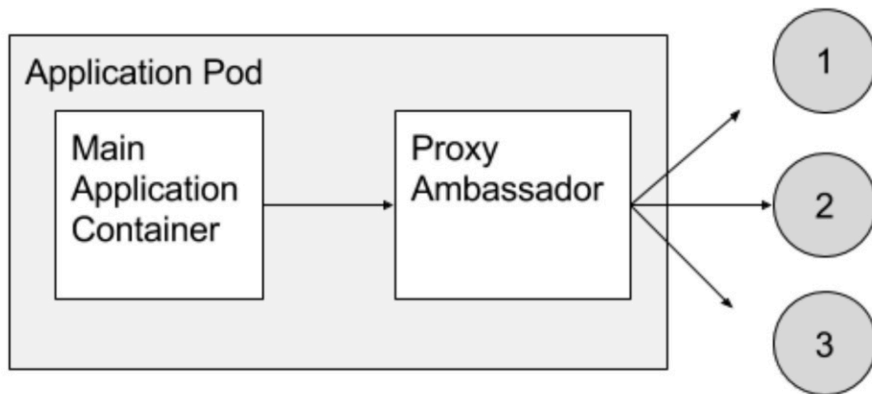
Multi-container Pods - Sidecar

- Additional container for support functionality allowing separation of concerns
- E.g. A log saving container that forwards logs to a central monitoring system



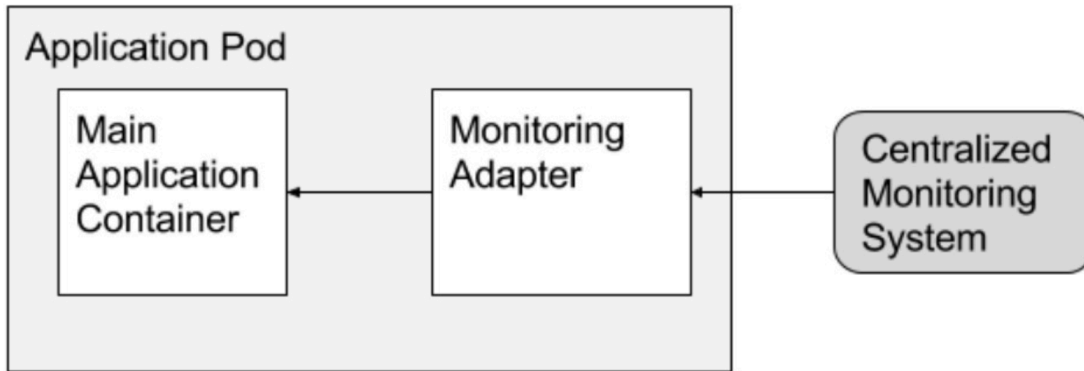
Multi-container Pods - Ambassador

- Proxy a local connection to the world.
- E.g. using a Redis server with one main write and multiple replicas for read. The ambassador proxy handles which redis instance to connect to, developers of main container don't need to add additional logic.



Multi-container Pods - Adapter

- Adapt or modify behaviour of main container
- E.g. Consider the task of monitoring N different applications, each with different monitoring data. Adapter can modify data to standardize for central monitoring system.



More about Pods 

Labels

- Useful for increasing search efficiency

```
apiVersion: v1
kind: Pod
metadata:
  name: my-production-label-pod
  labels:
    app: my-app
    environment: production
spec:
  containers:
    - name: nginx
      image: nginx
---
apiVersion: v1
kind: Pod
metadata:
  name: my-development-label-pod
  labels:
    app: my-app
    environment: development
spec:
  containers:
    - name: nginx
      image: nginx
```

```
# get pods with specific label filters
kubectl get pods -l app=my-app
kubectl get pods -l environment=production
kubectl get pods -l environment=development
kubectl get pods -l environment!=production
kubectl get pods -l 'environment in (development,production)'
kubectl get pods -l app=my-app,environment=production
```

Annotations

- Used to annotate objects
- Cannot be searched

```
apiVersion: v1
kind: Pod
metadata:
  name: my-annotation-pod
  annotations:
    owner: myaccount@k8sisgr8.com
    git-commit: bdab0c6
spec:
  containers:
  - name: nginx
    image: nginx
```

Container logs

```
# read logs from a specific container  
kubectl logs my-pod -c my-container
```

```
# output to a file  
kubectl logs my-pod -c my-container > my-container.log
```

```
# stream logs  
kubectl logs -f my-pod
```

Metrics

```
# see CPU and memory per pod
kubectl top pods

# specific pod
kubectl top pod my-resource-pod

# kube-system
kubectl top pods -n kube-system

# the nodes
kubectl top nodes
```

✨ This requires installation of metrics server (already on GKE)

Debugging

- Kubectl commands like `get`, `describe`, `edit` are your friend 😊

```
# list pods from all namespaces
kubectl get pods --all-namespaces

# get details about a pod
kubectl describe pod nginx -n nginx-ns

# edit a pod manifest
kubectl edit pod nginx -n nginx-ns
```


Cattle not pets 🐮

✨ Pods are intended to be immutable. Some fields can be updated (like `image`) but correct practice is to replace pods.





Tune in for more when we talk about Deployments. 📺

Exercise 1

1. Create a YAML file with a namespace definition for `sunnybikes` and apply it
2. Create a YAML file with a pod definition for the sunnybikes application. For this, use the docker image `pugillum/sunnybikeslite:stable`. Apply this manifest.
3. Apply the configurations and check the logs of the pods with:
`kubectl -n sunnybikes logs <podname>`
4. Port-forward using `kubectl port-forward pod/sunnybikes 8000:5000` and access <http://localhost:8000> in your browser. Now try <http://localhost:8000/docs>


Deployments

Pods

- Cannot self-heal 
- Don't scale 
- Updates and rollbacks -> hard 
- Enter Deployments 

Deployments

- Replicasets manage Pods
 - self-healing
 - scaling
- Deployments manage Replicasets
 - rollouts
 - rollbacks



Deployments - State

- Desired state - declared
- Observed state - monitored by Controllers (ReplicaSet and Deployment)
- Reconciliation - bring in sync

Deployments

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2 # tells deployment to run 2 pods matching the template
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.14.2
          ports:
            - containerPort: 80
```

```
# check deployment
kubectl get deployment
```

Rolling updates

- Update with zero downtime
- Various strategies, default RollingUpdate

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: rolling-deployment
spec:
  strategy:
    rollingUpdate:
      maxSurge: 3 # max pods that can be scheduled above desired number
      maxUnavailable: 2 # max pods that can be unavailable during update
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.7.1
          ports:
            - containerPort: 80
```

```
# set image version for specific deployment
kubectl set image deployment/rolling-deployment nginx=nginx:1.7.9

# View the rollout history of a deployment
kubectl rollout history deployment/rolling-deployment

# View the details of revision 2
kubectl rollout history deployment/rolling-deployment --revision=2
```


Deployment strategy - standard rolling

If `strategy` left out, the default looks like this:

```
spec:
  replicas: 3
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 25% # No more than 25% above desired state
      maxUnavailable: 25% # No more than 25% below desired state
```

Deployment strategy - recreate

- terminate all running instances then recreate with newer version.
- this means downtime

```
spec:  
  replicas: 3  
  strategy:  
    type: Recreate
```

Deployment strategy - ramped

- slow rollout
- for stateful apps that need to rebalance data

```
spec:  
  replicas: 3  
  strategy:  
    type: RollingUpdate  
    rollingUpdate:  
      maxSurge: 2 # No more than 2 pods above desired state  
      maxUnavailable: 0 # No pods below desired state
```

Deployment strategy - Blue/Green

- Have separate deployments for old and new (blue/green)
- Use a service to shift traffic from blue to green
- Can quickly roll back to blue if green has issues

Deployment strategy - Canary

- Similar to Blue/Green but only portion of users directed to new "Canary" deployment
- Rest of traffic shifted once real users have verified the new one

Rollback

```
# rollback to previous deployment  
kubectl rollout undo deployment/rolling-deployment  
  
# rollback to a specific version  
kubectl rollout undo deployment/rolling-deployment --to-revision=1
```

Exercise 2

1. Migrate the pod definitions for `sunnybikes` to a deployment (called `sunnybikes`)
2. The `sunnybikes` pods must be replicated at least 3 times
3. Make sure that when a new update of `sunnybikes` is deployed at least 1 pod is always available during upgrading and the number of pods can go up to a maximum of 6.

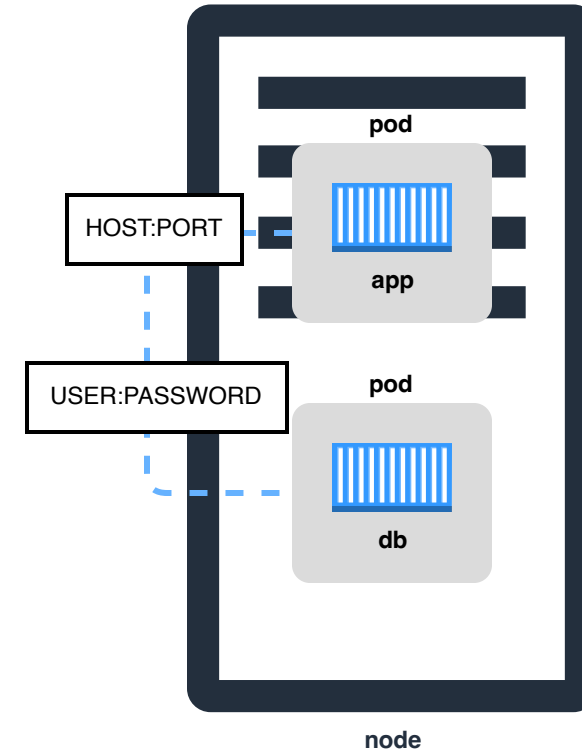
Configuration

Environment variables

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-env
spec:
  containers:
    - name: busybox
      image: busybox:stable
      command: ["sh", "-c", "echo $MESSAGE;"]
      env:
        - name: MESSAGE
          value: "Hello!"
```

Config Maps & Secrets

- configuration with ConfigMap
- Secret for secret data
- environment variables or properties file



ConfigMaps

- Sharing configuration between pods / containers
- Mountable as environment variables

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-configmap
data:
  # key value style
  message: hello
  name: John

  # file style
  app.cfg: |
    key1=value1
    key2=value2
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod-cf
spec:
  restartPolicy: Never
  containers:
    - name: busybox
      image: busybox:stable
      command: ["sh", "-c", "echo $MESSAGE $NAME;"]
      env:
        - name: MESSAGE
          valueFrom:
            configMapKeyRef:
              name: my-configmap
              key: message
        - name: NAME
          valueFrom:
            configMapKeyRef:
              name: my-configmap
              key: name
```


ConfigMaps

- Mountable as volume

```
apiVersion: v1
kind: Pod
metadata:
  name: my-configmap-volume-pod-1
spec:
  restartPolicy: Never
  containers:
  - name: myapp-container
    image: busybox
    command: ['sh', '-c', "echo $(cat /etc/config/hello)"]
    volumeMounts:
    - name: config-volume
      mountPath: /etc/config
  volumes:
  - name: config-volume
    configMap:
      name: my-config-map
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-configmap-volume-pod-2
spec:
  restartPolicy: Never
  containers:
  - name: busybox
    image: busybox:stable
    command: ["sh", "-c", "cat /config/app.cfg"]
    volumeMounts:
    - name: config
      mountPath: /config
      readOnly: true
  volumes:
  - name: config
    configMap:
      name: my-configmap
      items:
      - key: app.cfg
        path: app.cfg
```

Secrets

- Manage sensitive data like passwords, tokens etc.
-  beware of caveats - see [k8s documentation](#)

```
echo Secret Stuff! -n | base64  
U2VjcmV0IFN0dWZmMQo=
```

```
apiVersion: v1  
kind: Secret  
metadata:  
  name: my-secret  
data:  
  myKey1: U2VjcmV0IFN0dWZmMQo=  
stringData:  
  myKey2: myPassword
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: my-secret-pod  
spec:  
  containers:  
    - name: myapp-container  
      image: nginx:stable  
      env:  
        - name: MY_PASSWORD  
          valueFrom:  
            secretKeyRef:  
              name: my-secret  
              key: myKey
```

Exercise 3 - Part 1

1. Add an additional deployment `postgres` based on the image `postgres:11-alpine`, only one pod should run
2. Define an environment variable for the `sunnybikes` pods `PG_PORT` with value `"5432"`
3. Create a yaml file with a secret definition for the postgres password called `secret.yaml` and apply
4. Add the secret as environment variable to the sunny and postgres pods. SunnyBikes needs a `PG_PASSWORD` env variable and postgres needs a `POSTGRES_PASSWORD` env variable

Exercise 3 - Part 2

5. Create a yaml file with a config map (`configmap.yaml`) for the postgres init schema. Mount the init script in the postgres container in the folder `docker-entrypoint-initdb.d` as file `init-schema.sql`.

The schema is as follows:

```
CREATE TABLE public.bike_rides (  
  uuid UUID PRIMARY KEY,  
  name VARCHAR(80) NOT NULL,  
  location VARCHAR(80) NOT NULL,  
  created TIMESTAMPTZ NOT NULL  
)
```

Services


Services

- New pods == new IP addresses
- IP churn!
- Services provide stable network endpoints



Service

- Stable IP address
- Stable DNS
- Seamless connection across nodes
- Lifecycle detached from pods
- Internal and external
- Load balancing across pods
- Network policies



Services - implementing

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: nginx # determines which pods receive traffic
  ports:
    - protocol: TCP
      port: 8080 # port that services listens on
      targetPort: 80 # port on pod to route traffic to
```

```
apiVersion: v1
kind: Service
metadata:
  name: nodeport-service
spec:
  type: NodePort # a different type of service
  selector:
    app: nginx # determines which pods receive traffic
  ports:
    - protocol: TCP
      port: 8080
      targetPort: 80
      nodePort: 30080 # the port to map to on the node
```

Pods linked to these services

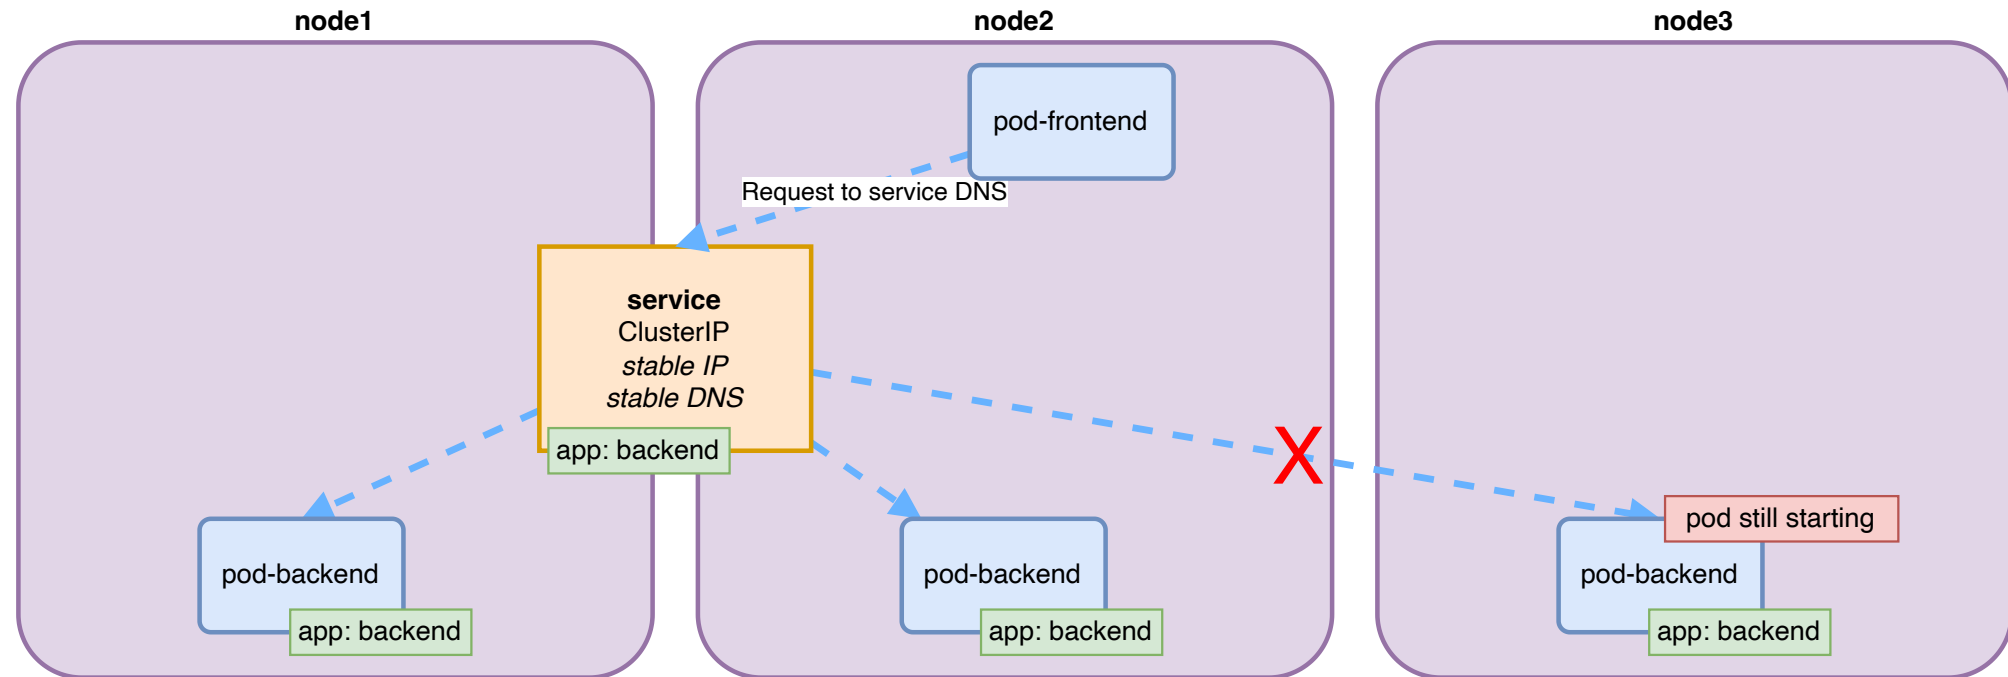
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deploy
spec:
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx # note this label compared to service selector!
  spec:
    containers:
      - name: nginx
        image: nginx:1.20.1
        ports:
          - containerPort: 80
```

```
kubectl get svc
kubectl get endpoints my-service
```

Services - ClusterIP


- The default type
- Exposes the Service on a cluster-internal IP
- Only reachable from within the cluster
- Has a DNS entry `<service-name>.<namespace>.svc.cluster.local`

Services - ClusterIP




Services - NodePort

- Exposes the Service on each Node's IP at a static port (the NodePort)
- Automatically creates a ClusterIP Service
- From external connect to `<NodeIP>:<NodePort>`



Services - LoadBalancer

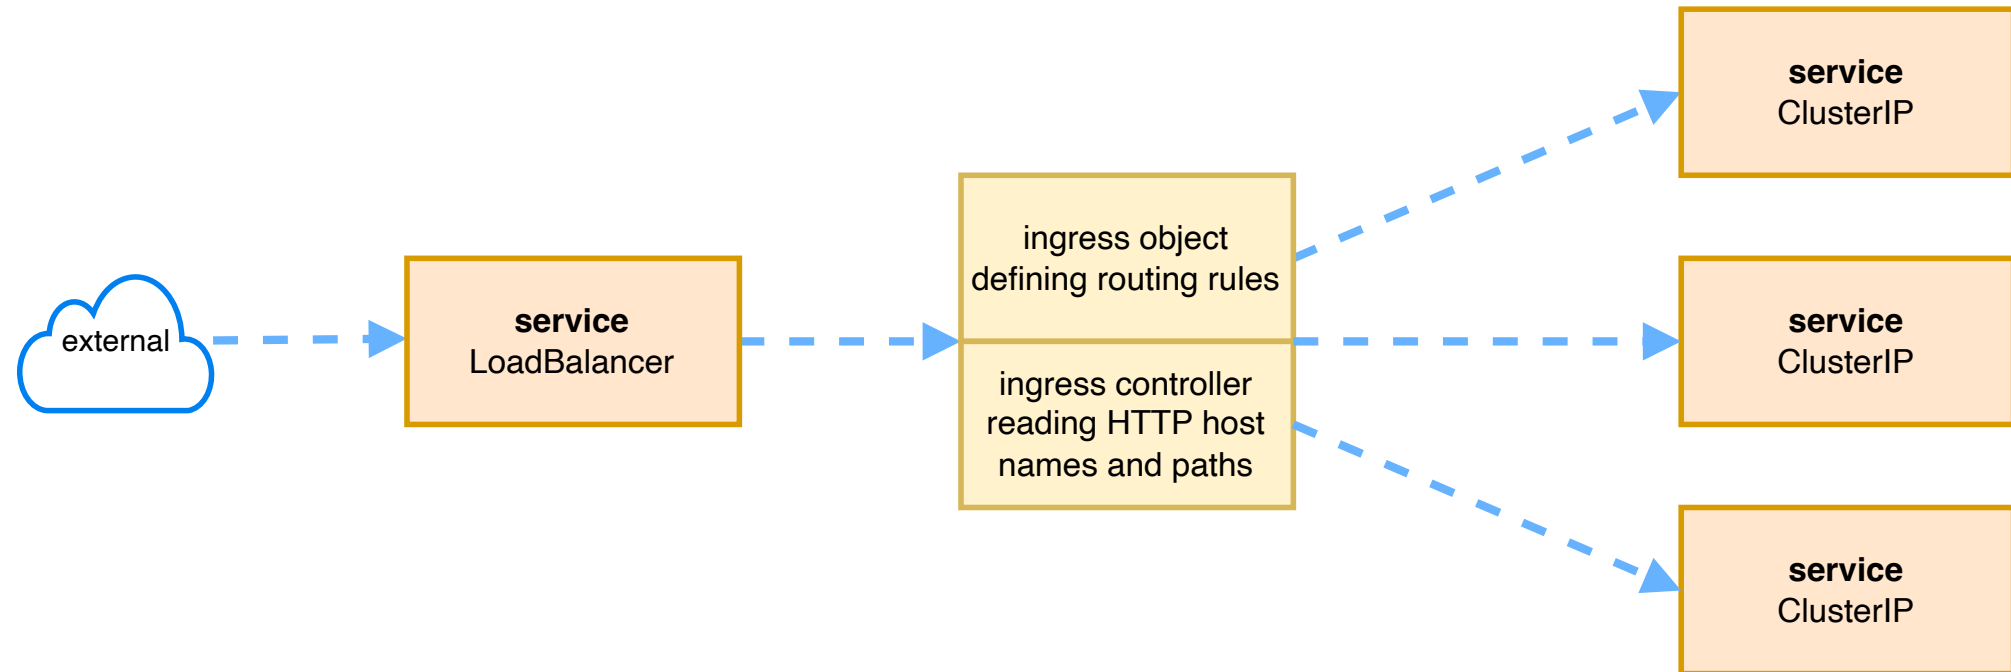
- Uses external load balancer (often from cloud provider)
- NodePort and ClusterIP Services, to which the external load balancer routes, are automatically created



Ingress

- can be used to expose *multiple* services
- not a Service type, but acts as entry point
- allows consolidating your routing rules into a single resource
- most powerful, but can also be the most complicated

Ingress



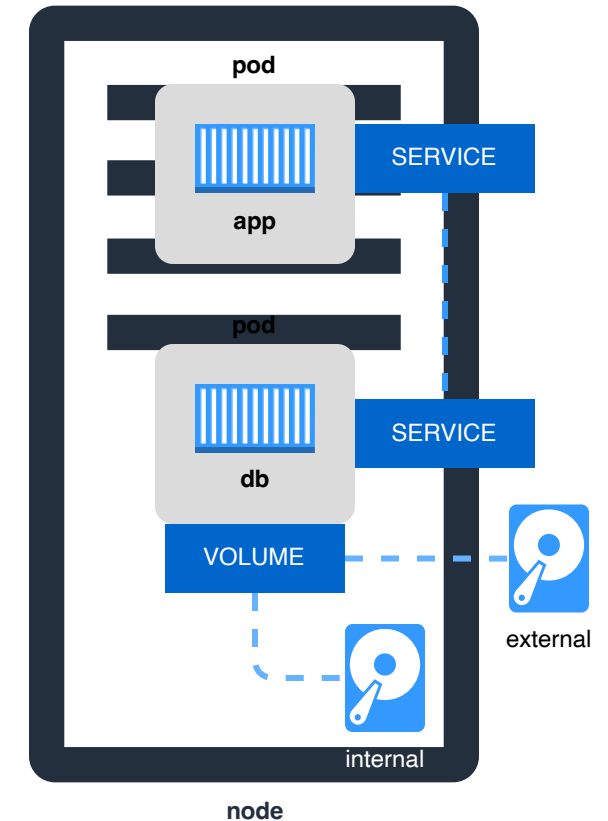
Exercise 4

1. Make Postgres available to the `sunnybikes` pods from only within the cluster
2. `sunnybikes` should be adjusted to use a new image `sunnybikes:stable`
3. `sunnybikes` takes 2 environment variables to configure the postgres host and port, `PG_HOST` - `<service name>.<namespace>.svc.cluster.local` and `PG_PORT` - `"5432"` respectively.
4. Make `sunnybikes` available on port 80 from the outside world

Storage

Volumes

- Pod disk storage is ephemeral 😞
- Volume - a directory that is accessible to all containers in a pod
- Ephemeral volumes - lifetime of the pod
- Persistent volumes - beyond pod lifetime
- Abstract a variety of different storage types - portability FTW!



Volumes in Pod definition

- Containers can share storage

```
apiVersion: v1
kind: Pod
metadata:
  name: volume-pod
spec:
  containers:
    - name: busybox1
      image: busybox:stable
      command: ["/bin/sh", "-c", 'echo "The writer wrote this!" > /output/data.txt; while true; do sleep 5; done']
      volumeMounts:
        - mountPath: /output
          name: my-volume
    - name: busybox2
      image: busybox:stable
      command: ["/bin/sh", "-c", "while true; do cat /input/data.txt; sleep 5; done"]
      volumeMounts:
        - mountPath: /input
          name: my-volume
  volumes:
    - name: my-volume
      emptyDir: {} # exists for lifetime of Pod
```

Kubernetes Storage Objects

PersistentVolume (PV)

Storage resources setup in the cluster

PersistentVolumeClaim (PVC)

Connect pods to storage via request

StorageClass (SC)

Dynamically provision PVs with specific characteristics

Container Storage Interface

- A standardized interface for storage integration
- Allows separate development of storage plugins



Persistent Volume

- Represent storage resources in the cluster
- Static PVs are provisioned by cluster admin
- Dynamic PVs are provisioned via StorageClass
- Independent of Pod lifecycle
- Local or Remote
- There are many types - each have a specific configuration
- Not linked to Namespace

Persistent Volume - types

- `local` – Data is stored on devices mounted locally to your cluster's Nodes.
- `hostPath` – Stores data within a named directory on a Node (designed for testing purposes)
- `nfs` – Used to access Network File System (NFS) mounts.
- `csi` – Allows integration with storage providers that support the Container Storage Interface (CSI) specification, such as the block storage services provided by cloud platforms.
- more...

Storage Class

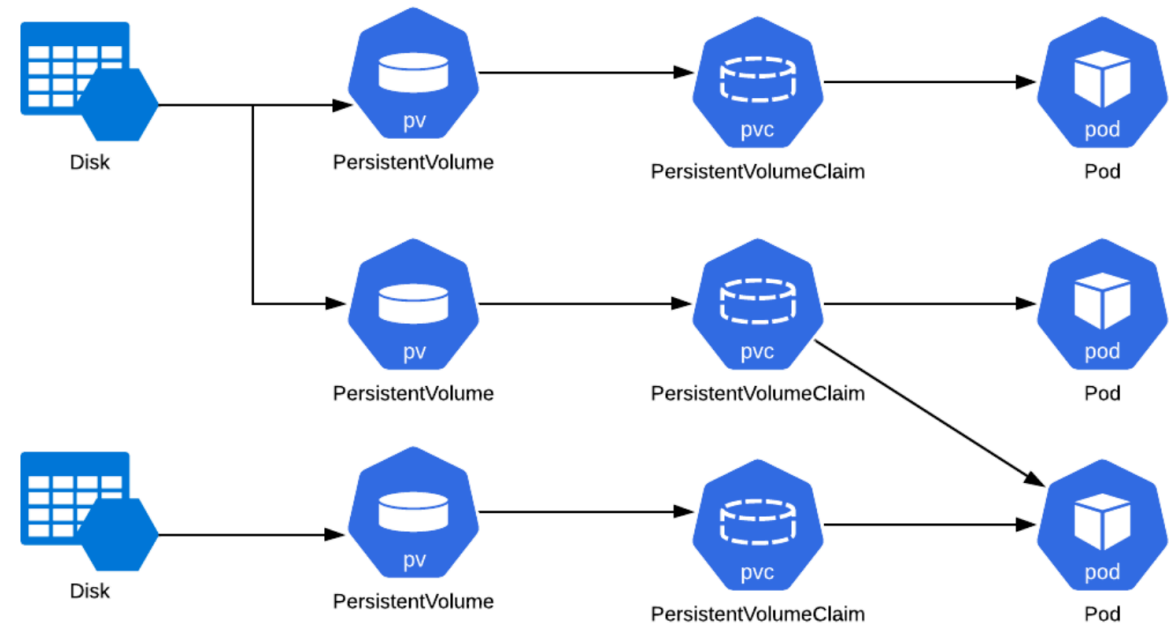
- Handle storage provisioning operations
- Managed by cluster admins
- There are built-in storage classes for each of the supported PV types
- Additional storage classes can be added by installing plugins

Persistent Volume Claim

- Defines a request for storage
- Includes details on type of storage needed (capacity, access mode and storage class)
- Automatically binds to available PersistentVolume that meets requirements
- Mounted in a pod like a volume (available to all containers)

Persistent storage

- A pod may have many PersistentVolumeClaims (PVC)
- A PVC may be used by many pods
- A PersistentVolume may only be claimed by one PVC
- A disk may have many PersistentVolumes requesting parts of storage



Persistent storage

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-pv
spec:
  storageClassName: manual # used for binding PVC
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteOnce # volume can be mounted as read-write by a single Node
  hostPath:
    path: "/mnt/data" # maps to this path on node
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-pvc
spec:
  storageClassName: manual # matches PV storage class
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 512Mi
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pvc-pod
spec:
  containers:
    - name: nginx
      image: nginx:stable
      ports:
        - containerPort: 80
      volumeMounts:
        - mountPath: "/usr/share/nginx/html"
          name: my-storage
  volumes:
    - name: my-storage
      persistentVolumeClaim:
        claimName: my-pvc
```

Exercise 5

If the postgres pod dies, all the data is lost. 😭 Make sure all the postgres data is persistent even throughout pod restarts

Hint: Add data via the Swagger interface of the API

Hint: Postgres stores its data in `/var/lib/postgresql/data`

Verify the volume is working.

Pods - Advanced

imagePullPolicy

- **IfNotPresent** - reduce network bandwidth
- **Always** - to ensure the latest version always - the default
- **Never** - image must already be on node, for security

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - name: mycontainer
    image: myimage:latest
    imagePullPolicy: IfNotPresent
```


Resource requirements

- `limits` : Restriction on resources for pods
- `requests` : Requests on resource for pods

```
apiVersion: v1
kind: Pod
metadata:
  name: my-resource-pod
spec:
  restartPolicy: Never
  containers:
    - name: nginx
      image: nginx:stable
      resources:
        requests:
          memory: "64Mi"
          cpu: "250m"
        limits:
          memory: "128Mi"
          cpu: "500m"
```

Job / CronJob

- Controller for short-lived pods (can be multiple)
- Job - one time or scheduled tasks
- CronJob - periodic tasks
- Tracks progress and ensures completion

Job / CronJob YAML

```
apiVersion: batch/v1
kind: Job
metadata:
  name: my-job
spec:
  template:
    spec:
      containers:
        - name: print
          image: busybox:stable
          command: ["echo", "This is a test!"]
          restartPolicy: Never
      backoffLimit: 4 # Optional: number of retries
      activeDeadlineSeconds: 10 # Optional: max time until job stopped
```

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: my-cronjob
spec:
  schedule: "*/1 * * * *" #
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox:stable
              command: ["echo", "This is a test!"]
              restartPolicy: OnFailure
```

Liveness probe

- Can customize how Kubernetes detects the status of **containers**
- Indicates if a container is running properly and governs when the cluster will automatically stop or restart the container.
- Health status via execution of command or HTTPGet

```
apiVersion: v1
kind: Pod
metadata:
  name: my-liveness-pod
spec:
  containers:
  - name: myapp-container:stable
    command: ["sh", "-c", "while true; do sleep 10; done"]
    livenessProbe:
      exec:
        command:
        - echo
        - testing
      initialDelaySeconds: 5
      periodSeconds: 5
```

Readiness probe

- Can customize how Kubernetes detects the status of **containers**
- Indicates whether a container is ready to service requests and governs whether requests will be forwarded to the pod.
- Ready to serve status via exec command or HTTPGet

```
apiVersion: v1
kind: Pod
metadata:
  name: my-readiness-pod
spec:
  containers:
  - name: nginx
    image: nginx:1.20.1
    readinessProbe:
      httpGet:
        path: /
        port: 80
      initialDelaySeconds: 5
      periodSeconds: 5
```

Exercise 6

1. Add a liveness probe to check the `healthz` endpoint of the sunny bikes API.
2. Add a readiness probe to check that Postgres is ready by calling the command `pg_isready -U postgres` with a delay of 10 seconds, repeating at 10 second intervals.

Extra topics

Roles Based Access Control (RBAC)

- Role: define what is allowed
- RoleBinding: tie Roles to users or service accounts
- Varies per cloud provider

```
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: pod-reader
  namespace: default
rules:
- apiGroups: [""] # don't need special API group
  resources: ["pods"] # type of k8s object
  verbs: ["get", "watch", "list"] # what we can do
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: pod-reader-binding
subjects: # provides list of accounts that this role binding is addressing
- kind: User
  name: imauser@k8sworkshop.com
  namespace: default
roleRef: # which role we're binding subjects to
  kind: Role
  name: pod-reader
  apiGroup: rbac.authorization.k8s.io
```


Security

- security flaws sometimes exist
- there are practices to be safe
- run as non-root as much as possible
- security is hard
- align with IT



Non-root - Python Dockerfile

- `USER` configures user for subsequent `RUN`, `CMD` and `ENTRYPOINT`

```
FROM python:3.12-slim

COPY requirements.txt requirements.txt
RUN pip install --user -r requirements.txt

RUN groupadd -r workers && useradd -r -g workers worker
USER worker
WORKDIR /home/worker

ENV PATH="/home/worker/.local/bin:${PATH}"

COPY --chown=worker:workers . .

ENTRYPOINT ["python"]
```

Pod Security context


`runAsNonRoot` - simple and safe

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
  - name: myapp-container
    image: busybox:stable
    securityContext:
      runAsNonRoot: True
```


```
apiVersion: v1
kind: Pod
metadata:
  name: security-context-demo
spec:
  securityContext:
    runAsUser: 1000
    runAsGroup: 3000
    fsGroup: 2000
  volumes:
  - name: sec-ctx-vol
    emptyDir: {}
  containers:
  - name: sec-ctx-demo
    image: busybox:1.28
    command: [ "sh", "-c", "sleep 1h" ]
    volumeMounts:
    - name: sec-ctx-vol
      mountPath: /data/demo
    securityContext:
      allowPrivilegeEscalation: false
```

StatefulSet

- for stateFUL
- synchronizes reads/writes
- ensures data consistency
- ⚠ not advised - external storage better



Databricks Containers

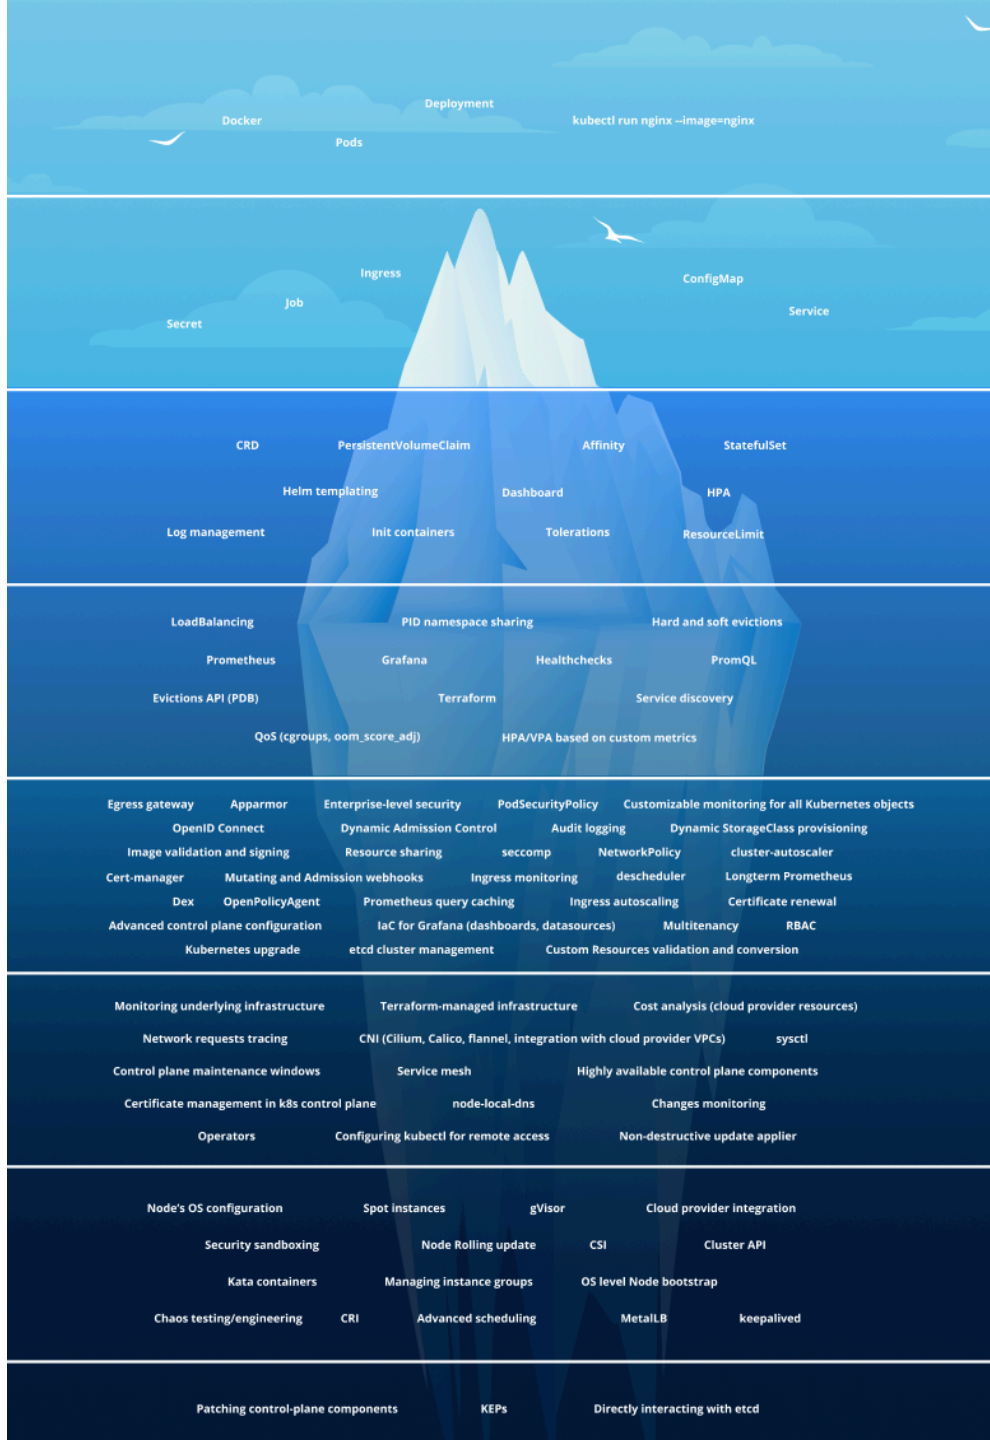
- Databricks Container Service
- Can use Databricks base images or from scratch
- Only LTS images patched regularly
- Use with caution 

Helm

- A package manager for Kubernetes
- Simplifies deployment and management of applications on a cluster
- Combine multiple manifests into a single chart
- Templating
- Version control
- Rollout/rollback easily
- Community maintained charts

Kustomize

- Customize YAML files without the need for templates
- Define base templates in standard manifest YAML
- Use patching to adjust settings per environment
- Manage objects with a `kustomization` file
- Installed with `kubectl`



Still more to learn 🤖



Thanks for listening! 😊